

STAT 315 Notes Fall 2026

Colorado State University

Preliminaries

Getting Started With R

Installing R

R is the base version of a free, open-source statistical computing program. RStudio is a nice graphic user interface (GUI) that works in conjunction with R. We will be using RStudio this semester as it is easier to learn how to use and has many nice shortcuts that aren't available in the standard R download.

Note: R and R Studio can not be installed on Chromebooks or other tablets/mini-computers. If you only have a tablet or a mini-computer, the free version of RStudio Cloud should work: R Studio Cloud. The free version allows 25 hours of use per month. I think this is likely enough for this course, but be sure to log out!

To begin using RStudio, you will need to install “R” first and then install “RStudio” on your computer.

Step 1: Download R

- Visit <https://www.r-project.org/>
- Click *CRAN* under *Download*
- Select any of the mirrors
- Click the appropriate link for your type of system (Mac, Windows, Linux)
- Download R on this next page.
 - For Windows, this will say *install R for the first time*
 - For Mac, this will be under *Latest release* and will be something like *R-4.1.0.pkg* – the numbers may differ depending on the most recent version
- Install R on your computer

Step 2: Download RStudio

- Visit <https://posit.co/downloads>
- Click to download
- Click the appropriate link for your type of system (Mac, Windows, Linux)
- Install RStudio on your computer

Step 3: Verify RStudio is working

- Open RStudio
- Let's enter a small dataset and calculate the average to make sure everything is working correctly.
- In the console, type in the following dataset of five homework scores:

```
# Assignment using =  
x = c(90,75,100,85,90)  
  
# Alternative assignment using <-  
x <- c(90,75,100,85,90)  
  
# More on the differences between = and <- at:  
# https://www.r-bloggers.com/2014/01/difference-between-assignment-operators-in-r/
```

In the console, calculate the average of the five homework scores:

```
mean(x)
```

```
## [1] 88
```

Did you find the average of the five homework scores was 88? If so, you should be set up correctly!

Basic Concepts of Statistics

In this chapter, we'll introduce some of the basic concepts and definitions that are fundamental to the study of Probability and Statistics. In particular, sampling and sample statistics will be examined.

Populations, Samples, and Processes

A **population** is a well-defined complete collection of objects.


A **sample** is a subset of the population.

A **population parameter** is a numerical value that describes a characteristic of a population.

A **sample statistic** is a numerical value that describes a characteristic of a sample.

A student is interested in the average grade point average (GPA) of CSU students. She collects the GPAs of eight friends and finds an average of 2.87. For this example, determine the population, sample, parameter, and statistic.

Illustration of Population and Sample



Methods of Collecting a Sample

Simple Random Sampling (SRS) is a basic sampling method used in statistics where every individual or element in a population has an equal chance of being selected. This method ensures that the sample is unbiased.

Suppose we want to collect a simple random sample of CSU student GPAs. How might we do this?

Convenience Sampling is a sampling method where the sample is taken from a group of people that are easy to reach or readily available. This method relies on selecting individuals who are most conveniently accessible to the researcher, rather than using a random or systematic approach.

Suppose we want to collect a convenience sample of CSU student GPAs. How might we do this?

Stratified Sampling is a probability sampling method where the population is divided into distinct subgroups or “strata” that share similar characteristics, and a random sample is then taken from each stratum. This approach ensures that each subgroup is adequately represented in the final sample, leading to more precise and reliable estimates.

Suppose we want to collect poll data for the presidential election for 250 voters. Further, suppose a recent poll found that 36% of voters are registered with the Democratic Party, 31% of voters are registered with the Republican Party, and 33% are not registered for either major party. How might we collect a stratified sample and what’s the advantage compared to SRS?

Note: While there are many other more complex techniques to sample, we’ll generally assume that any sample collected in this class is a simple random sample.

Data can be either collected in an *observational study* or in an *experiment*.

In an ***observational study***, data is collected without manipulating the environment or influencing how the data naturally occur.

How might one collect data to investigate a connection between smoking and cancer?

In an ***experiment***, participants are assigned to different treatment groups to determine if the treatment causes changes in the outcome. A key component is the ***control group***, where participants receive no treatment or a standard one, serving as a baseline for comparison. This setup helps establish causality by isolating the effect of the treatment from other variables.

Why might it be difficult to conduct a controlled experiment investigating smoking and cancer?

Suppose we want to set up a controlled experiment to investigate the efficacy of a new blood pressure reduction drug. We begin by randomly sampling 100 people with high blood pressure. How would we set up a controlled experiment?

In addition to the sampling method, bias may still occur if participants in an experiment are aware of the treatment they are receiving. This awareness can influence their behavior or responses, potentially skewing the results.

20 patients are being treated for depression, half are given an antidepressant and half are given a sugar pill (placebo). Suppose the patients know which pill they have been given. How might this cause the results to be biased?

Blinded Experiments

A ***blinded experiment*** is an experiment in which participants do not know which treatment they are receiving.

A ***double-blinded experiment*** is an experiment in which both the participants and the researchers do not know which treatment is being administered during the experiment.

A ***triple-blinded experiment*** is an experiment in which the participants, the researchers, and the analysts evaluating the data are all unaware of which treatment is being administered. This approach further reduces the potential for bias at all stages of the experiment.

Types of Data

We will break data into two categories: **quantitative** and **qualitative**.

Quantitative data is numeric data or numbers and can be broken into two further categories, ***discrete*** and ***continuous***.

Discrete data is quantitative data with a finite or countably infinite number of values.

Continuous data is quantitative data with an uncountably infinite number of values or data taken from an interval.

Give examples of discrete and continuous data.

Qualitative data refers to names, categories, or descriptions and can also be broken down into two further categories, ***nominal*** and ***ordinal***.

Nominal data is qualitative data with no natural ordering.

Ordinal data is qualitative data with a natural ordering.

Give examples of nominal and ordinal data.

Summarizing Qualitative Data

While we will mostly focus on quantitative data in this class, there are a few simple methods for describing qualitative data that you should be aware of.

Let's collect some data on eye color from members of the class and then explore some statistical and graphical summaries.

Data:

*Frequency Table**Bar Graph*

Proportions: $p = \frac{\# \text{ in category}}{\# \text{ total}}$

What proportion are brown?

Percentages: $P\% = p \cdot 100$

What percentage are blue?

Summary Statistics for Quantitative Data

For the following exercises, we will use the following data set. The heights of five students are measured to be (in cm): 183, 165, 165, 175, 187.

Notation: $\sum_{i=1}^n x_i = x_1 + x_2 + \dots + x_n$

What is $\sum x_i$, the sum of all heights in the sample?

There are many ways to measure the “center” of our data set. We will look at mean, median, and mode.

The **sample mean** (or **sample average**) $\bar{x}$ is given by: $\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$.

What is the sample mean of the original heights data set, $\bar{x}$?

The **sample median** $\tilde{x}$ is obtained by first ordering the n observations from smallest to largest and then calculated by the following:

$$\tilde{x} = \begin{cases} \text{the single middle value if } n \text{ is odd,} & \tilde{x} = x_{(n+1)/2} \\ \text{the average of the two middle values if } n \text{ is even,} & \tilde{x} = (x_{n/2} + x_{n/2+1})/2 \end{cases}$$

where x_i is the i^{th} smallest value in the data set.

What is the sample median of the heights data set, $\tilde{x}$?

If we measure an additional person at 174cm, what is the new median of the data set?

The ***sample mode*** is the most commonly occurring data point.

What is the sample mode of the heights data set?

Measures of Variability

There are a few different ways to describe the variability of a data set, including range, variance, standard deviation, and interquartile range.

The ***range*** of a data set is the difference between the largest and the smallest values.

What is the range of the heights data set?

Typically, the most common method for measuring the spread of a data set is variance or standard deviation.

The **sample variance**, denoted by s^2 , is given by:

$$s^2 = \sum_{i=1}^n \frac{(x_i - \bar{x})^2}{n-1}$$

The **sample standard deviation**, denoted by s , is given by:

$$s = \sqrt{s^2} = \sqrt{\sum_{i=1}^n \frac{(x_i - \bar{x})^2}{n-1}}$$

Using the heights data set, calculate the sample variance and standard deviation.

i	1	2	3	4	5	Σ
x_i						
$x_i - \bar{x}$						
$(x_i - \bar{x})^2$						

The **z-score** measures the number of standard deviations an observation is from the mean: $z_i = \frac{x_i - \bar{x}}{s}$

A **percentile** is a measure of relative standing. The p^{th} percentile is the number where at least $p\%$ of the data values are less than or equal to this number.

Special percentiles:

1. 25^{th} percentile = 1^{st} quartile = Q_1
2. 50^{th} percentile = 2^{nd} quartile = Q_2 = median
3. 75^{th} percentile = 3^{rd} quartile = Q_3

Computing Quartiles

1. Order the list of data values from smallest to largest
2. Compute the median. This is Q_2 .
3. Split the data in half at the median. If there are an odd number of data values, the median is included in neither half.
4. Q_1 is computed by finding the median of the lower half.
5. Q_3 is computed by finding the median of the upper half.

Calculate Q_1 and Q_3 of the heights data set.

The ***interquartile range (IQR)*** is a measure of statistical dispersion, representing the difference between the first quartile (Q_1) and the third quartile (Q_3). It captures the range of the middle 50% of the data, indicating the spread of the central portion of the dataset. It is given by: $Q_3 - Q_1$.

Calculate IQR for the heights data set.

An ***outlier*** is a data point that lies far outside of the normal range of the rest of the data.

Rule of thumb: Data points smaller than $Q_1 - 1.5 \cdot IQR$ or larger than $Q_3 + 1.5 \cdot IQR$ are outliers.

The ***five-number summary*** is a statistical description of a dataset, consisting of the minimum, first quartile (Q_1), median (Q_2), third quartile (Q_3), and maximum. It provides a quick overview of the data's distribution, highlighting its center and spread.

The five numbers are: minimum, Q_1 , median, Q_3 , maximum.

What is the five number summary for the heights data set?

Visualizing Quantitative Data


A **histogram** visualizes the distribution of a dataset by grouping data into bins and plotting the frequency of data points in each bin. The x-axis represents the data range, and the y-axis shows the frequency of occurrences. Histograms are useful for understanding the shape, central tendency, and spread of data.

14 students were randomly sampled and were asked how many credits they are taking this semester.

Data: 12, 14, 15, 12, 9, 18, 22, 15, 16, 14, 14, 12, 3, 17

Class	Frequency	Relative Frequency	Percentage Frequency
[0, 3)			
[3, 6)			
[6, 9)			
[9, 12)			
[12, 15)			
[15, 18)			
[18, 21)			
[21, 24)			
Total			


Histogram




```
## Loading required package: gridExtra
##
## Attaching package: 'gridExtra'
## The following object is masked from 'package:dplyr':
##
##      combine
```

Histogram shapes


1. Unimodal




2. Bimodal





3. Bell-shaped




4. Positively (right) skewed - tail is on the right



5. Negatively (left) skewed - tail is on the left**6. Symmetric**


Box Plots

A **box plot**, or box-and-whisker plot, visualizes data distribution using a five-number summary: minimum, first quartile, median, third quartile, and maximum [more on these later]. The box shows the interquartile range (IQR), and the whiskers extend to values within 1.5 times the IQR. Outliers are plotted as individual points. Box plots are useful for comparing distributions and spotting outliers.




Here are some examples of box plots.

Distribution of IQ Scores



Distribution of Minutes Late



Notation for Population Parameters and Sample Statistics

Some care must be given when reporting a mean, variance, or standard deviation, as it is more important to distinguish between population and sample.


	Mean	Median	Variance	Standard Deviation
Population	μ	$\tilde{\mu}$	σ^2	σ
Sample	$\bar{x}$	$\tilde{x}$	s^2	s

Descriptive Statistics vs. Inferential Statistics

After a sample is taken, there are a variety of ways to describe the sample. This is the realm of ***Descriptive Statistics***. We can describe our sample with numbers (statistics) or pictorially.

Later in the semester, we will examine ***Inferential Statistics***. In Inferential Statistics, our goal is to use characteristics of a sample to say something about the full population (see Inferential Statistics).

Connection between Probability and Statistics



If we know the characteristics of a population, the study of Probability will help us determine the likelihood that a sample will have certain characteristics.

Likewise, Statistics can potentially help us use the characteristics of a sample to say something about the population.

R Companion for Chapter 1

Example 1: Let's manually enter the heights data set (183, 170, 160, 175, 187) from chapter 1 and the following R commands will do the calculations we completed by hand.

Note: Lines of code that begin with `#` are comments and are not executed by the computer. Lines of code that begin with `#` denote the output.

```
# manually enter your data by using the "c" command  
heights = c(183,165,165,175,187)
```

```
# calculate sample statistics on data set  
mean(heights)
```

```
## [1] 175
```

```
median(heights)
```

```
## [1] 175
```

```
var(heights)
```


```
## [1] 102
```

```
sd(heights)
```

```
## [1] 10.0995
```

We can also create a histogram in R. Note that we can specify the breaks for the bins. In this case, the breaks are defined at 160, 170, 180, 190.

```
hist(heights,breaks=c(160,170,180,190))
```



Example 2: Let's create a randomized dataset of eye colors and analyze the data.

Note: When we generate random data in this class, we will use the **set.seed** function, so that we all generate the same random data. If we don't use the **set.seed** command before generating random data, we will all get different random data.

Here, we generate 47 random eye colors and take a look at the first few observations.

```
set.seed(2020)
eye.colors = sample(c("amber", "blue", "brown", "gray", "green",
                     "hazel", "red"), 47, replace = T)
```

```
# the "head" function will give the first few data values
# or the "header" of the dataset
head(eye.colors)
```


```
## [1] "gray" "gray" "red" "hazel" "amber" "amber"
```

We can create a frequency table of these values and create a barplot.

```
table(eye.colors)
```

```
## eye.colors
## amber blue brown gray green hazel red
##      3   10    3   10    7    8    6
```

```
barplot(table(eye.colors), xlab = "eye color",
         ylab = "frequency", main = "Bar plot of eye color occurrence")
```



We can also calculate proportion and percentage tables.

```
table(eye.colors) # shows how many of each value exist
```

```
## eye.colors
##  amber  blue brown  gray green hazel   red
##    3    10    3    10    7    8    6
```

```
length(eye.colors) # tells how many data points there are
```

```
## [1] 47
```

```
table(eye.colors)/length(eye.colors) # proportion
```

```
## eye.colors
##      amber      blue      brown      gray      green      hazel      red
## 0.06382979 0.21276596 0.06382979 0.21276596 0.14893617 0.17021277 0.12765957
```

```
table(eye.colors)/length(eye.colors)*100 # percent
```

```
## eye.colors
##      amber      blue      brown      gray      green      hazel      red
##  6.382979 21.276596  6.382979 21.276596 14.893617 17.021277 12.765957
```

Probability

In this chapter, we'll delve into Probability, the branch of mathematics that deals with the study of randomness and uncertainty, helping us understand and quantify the likelihood of different outcomes.

Sample Space and Events

An *experiment* is any activity or process whose outcome is subject to uncertainty.

The *sample space* of an experiment, denoted by Ω or $\mathcal{S}$, is the set of all possible outcomes of that experiment.

An *event* is any collection (subset) of outcomes contained in the sample space Ω .

For the following examples, define the sample space and give an example of an event. Code for running the experiment in R is given.

Experiment 1: Roll a fair six-sided die.

```
# roll one fair six-sided die in R
sample(x=1:6,size=1)
```

```
## [1] 5
```

Experiment 2: Roll two fair six-sided dice.

```
# roll two fair six-sided dice in R
sample(x=1:6,size=2)
```

```
## [1] 1 3
```

Experiment 3: Randomly select a real number in the closed interval from 0 to 1.

```
# generate one random number between 0 and 1
runif(n=1)
```

```
## [1] 0.6137196
```


Set Theory

Note: For the following examples, suppose a fair six-sided die is rolled, so $\Omega = \{1, 2, 3, 4, 5, 6\}$.

Define the following events: $A = \text{roll an even} = \{2, 4, 6\}$, $B = \text{roll a 3 or higher} = \{3, 4, 5, 6\}$, $C = \text{roll a prime number} = \{2, 3, 5\}$.


A **complement** of an event A , denoted by A^c or A' , is the set of all outcomes in Ω that are not contained in A .

Illustration




The **union** of two events A and B , denoted by $A \cup B$ and read “A or B”, is the event consisting of all outcomes that are either in A or B or in both.

Illustration




The **intersection** of two events A and B , denoted by $A \cap B$ and read “A and B”, is the event consisting of all outcomes that are in both A and B .

Illustration



The **set difference** of two events A and B , denoted by $A \setminus B$ and read “difference of A and B”, is the event consisting of all outcomes that are in A but not in B .

Illustration



The **null or empty set** is the set of no events (or elements), denoted $\emptyset$.


When $A \cap B = \emptyset$, A and B are said to be **mutually exclusive (or disjoint)** events.

Suppose a fair six-sided die is rolled, so $\Omega = \{1, 2, 3, 4, 5, 6\}$.


Define the following events: $A = \text{roll an even} = \{2, 4, 6\}$, $B = \text{roll a 3 or higher} = \{3, 4, 5, 6\}$, $C = \text{roll a prime number} = \{2, 3, 5\}$.

Determine the following sets and fill in the Venn diagrams representing these sets.


Draw and determine $A \cup B^c$.




Draw and determine $A \cap B \cap C^c$.



Draw and determine $[A \cup B] \cap C$.



Draw and determine $A \cap B \cap C$.



Axioms and Rules of Probability

The probability of an event, A , is denoted by $P(A)$ and represents the likelihood or chance that the event A will occur. Note that $0 \leq P(A) \leq 1$.

Axioms of Probability

The axioms of probability were introduced by the mathematician Andrey Kolmogorov in 1933. These axioms form the foundation of probability theory, providing a rigorous, axiomatic framework for quantifying uncertainty. In mathematics, an axiomatic approach means that all rules and theorems are built from a set of fundamental principles, or axioms.

Axiom 1: For any event A , $P(A) \geq 0$

Axiom 2: $P(\Omega) = 1$

Axiom 3: If $A_1, A_2, A_3, \dots$ is an infinite collection of disjoint events, then

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = P(A_1 \cup A_2 \cup A_3 \cup \dots) = \sum_{i=1}^{\infty} P(A_i)$$

Probability Rules

The following can be derived from the Axioms of Probability.

1. $P(\emptyset) = 0$

2. For any event A , $P(A) + P(A^c) = 1$

3. $P(A^c) = 1 - P(A)$

4. For any event A , $0 \leq P(A) \leq 1$

5. For any two events A and B , $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

6. For any three events A , B , and C ,

$$P(A \cup B \cup C) = P(A) + P(B) + P(C) - P(A \cap B) - P(A \cap C) - P(B \cap C) + P(A \cap B \cap C).$$

7. $[A \cup B]^C = A^C \cap B^C$

8. $[A \cap B]^C = A^C \cup B^C$

Suppose our class has the following breakdown:

70% Math students

20% seniors

30% Born in Denver

10% Math students and seniors

15% Math students and born in Denver

10% Seniors and Born in Denver

5% Math students, seniors, and born in Denver

Define events as follows: A = Math student, B = senior, C = Born in Denver.

Suppose a random student from the class is selected. Calculate the following probabilities.

(a) $P(A \cup B)$

(b) $P(C \setminus B)$

(c) $P([A \cap C]^c)$

(d) $P(A \cup B \cup C)$

Counting Techniques

Equally Likely Outcomes

If $\Omega = \{x_1, x_2, \dots, x_N\}$ and all outcomes are equally likely, then $P(A) = |A|/N$, where $|A|$ is the number of outcomes contained in the event A .

You roll a fair six-sided die. Calculate $P(\text{even})$.

(Product Rule for Ordered Pairs) If the first element of an ordered pair can be selected in n_1 ways and for each of these n_1 ways the second element of the pair can be selected in n_2 ways, then the number of pairs is $n_1 \cdot n_2$.

You flip two fair coins. Calculate the probability that all are heads or all are tails.

(Generalized Product Rule) Suppose a set consists of k elements (k -tuples) and that there are n_1 possible choices for the first element, n_2 possible choices for the second element, $\dots$, and n_k possible choices for the k^{th} element, then there are $n_1 \cdot n_2 \cdot \dots \cdot n_k$ possible k -tuples.

Quartiles is a word game where one of the goals is to find all five words that contain four tiles each. Suppose we have found 3 of the quartiles. How many total attempts are there to find the fourth quartile?



The factorial function, $n!$, is important in counting methods and probability. It is the product of all positive integers less than or equal to n , that is, $n! = n \cdot (n-1) \cdot (n-2) \cdot \dots \cdot 2 \cdot 1$ for $n \in \mathbb{Z}^+$.

Note: $0! = 1$

An ordered subset is called a **permutation**. The number of permutations of size k that can be formed from the n elements in a set will be denoted by $P_{n,k}$.

$$P_{n,k} = \frac{n!}{(n-k)!}$$

An unordered subset is called a **combination**. The number of combinations of size k that can be formed from the n elements in a set will be denoted by $\binom{n}{k}$ or $C_{n,k}$.

$$\binom{n}{k} = C_{n,k} = \frac{n!}{k! \cdot (n-k)!}$$

Suppose there are five competitors in an Olympic competition. A gold, a silver, and a bronze medal are awarded for the top three competitors. Are combinations or permutations more appropriate for this situation? How many possible ways can these three awards be given out?

Suppose there are five competitors in an Olympic competition and only three can advance to the final competition. Are combinations or permutations more appropriate for this situation? How many possible ways can you choose three people from a group of five?

Conditional Probability

Conditional probability is the likelihood of an event occurring given that another event has already occurred. It refines our understanding of probability by taking into account additional information that might affect the outcome. Understanding conditional probabilities is crucial in many areas of statistics and decision-making, as it allows us to update our predictions and assessments based on new evidence.

For any two events A and B with $P(B) > 0$, the conditional probability of A given B is defined by:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

The following is a contingency table from a survey on congressional approval.

Political Affiliation	Approves of Congress	Disapproves of Congress	Total
Republican	100	20	120
Democrat	20	80	100
Independent	40	40	80
Total	160	140	300

- (a) What is the probability that a randomly selected individual from the survey approves of Congress?

- (b) What is the probability that a randomly selected individual approves of Congress given that they are Democrat?

- (c) Given that the randomly selected person approves of the Congress, what is the probability that they are Republican?


Multiplication Rule: $P(A \cap B) = P(A|B) \cdot P(B)$

Suppose that people with a genetic variation have a 50% chance of getting a rare disease. The probability of having the genetic variation is 1%. What is the probability that a randomly chosen person has the genetic variation and also has the disease?

Law of Total Probability: Let $A_1, A_2, \dots, A_k$ be mutually exclusive and exhaustive events (i.e., the events form a partition). Then for any event B ,

$$P(B) = P(B|A_1) \cdot P(A_1) + P(B|A_2) \cdot P(A_2) + \dots + P(B|A_k) \cdot P(A_k)$$

Illustration of a Partition



Box #1 has ten computer chips (five working and five broken). Box #2 has five (one working and four broken). Suppose we randomly select one of the two boxes and then randomly select a chip. What is the probability that the chip is broken?

Bayes' Theorem: Suppose that $A_1 \cap A_2 = \emptyset$ and $A_1 \cup A_2 = \Omega$, then if $P(B) > 0$,

$$P(A_1|B) = \frac{P(B|A_1)P(A_1)}{P(B|A_1)P(A_1) + P(B|A_2)P(A_2)}$$

Let $A_1, A_2, \dots, A_k$ be a collection of mutually exclusive and exhaustive events, then if $P(B) > 0$,

$$P(A_j|B) = \frac{P(A_j \cap B)}{P(B)} = \frac{P(B|A_j) \cdot P(A_j)}{\sum_{i=1}^k P(B|A_i) \cdot P(A_i)}$$

Suppose a test to detect a rare disease is positive 95% of the time for a person with a disease and is positive 1% of the time for a person without the disease. Further, assume that 0.5% of the population has the disease. Given that a person tests positive for a disease, what is the probability that they actually has the disease?

Independence

Recall:

When $A \cap B = \emptyset$, A and B are said to be *mutually exclusive (or disjoint)* events.

Two events A and B are *independent* if $P(A|B) = P(A)$ and *dependent* otherwise.

Also, A and B are independent if and only if $P(A \cap B) = P(A) \cdot P(B)$.

Recall from the previous example that a test to detect a rare disease is positive 95% of the time for a person with a disease and is positive 1% of the time for a person without the disease. Further, assume that 0.5% of the population has the disease. Let D be the event that one has the disease and let P be the event that one has a positive test.

(a) Are D and P mutually exclusive?

(b) Are D and P independent?

Objective vs. Subjective Probability

There are two different interpretations of Probability. Both follow the Axioms of Probability.

Objective Probability refers to the probability of an event as determined by the long-term frequency of its occurrence, based on repeated trials or observations. It is rooted in empirical data and reflects the likelihood of an event as the proportion of times it occurs in a large number of trials.

I flip a coin 100 times and it lands on head 53 times. I estimate the probability of flipping a heads to be $P_{100}(\text{head}) = \frac{53}{100} = 0.53$. As we flip the coin more and more, we would expect that $P_n(\text{head}) \rightarrow 0.5$ as $n \rightarrow \infty$ if the coin is actually a fair coin.

Subjective Probability refers to the probability of an event as determined by an individual's personal judgment or belief, rather than objective data. It reflects the degree of confidence one has in the occurrence of an event. A well-known approach to subjective probability is Bayesian probability, which updates these beliefs based on new evidence.

I am given a coin which I have no reason to believe it is unfair (i.e., $P(\text{head}) = 0.5$). I flip the coin 100 times and 3 land on heads. Since this is very unlikely to happen if the coin is fair, I update my estimate to be $P(\text{head}) = 0.1$ (i.e., somewhere in between $3/100 = 0.03$ and $1/2 = 0.5$).

R Companion for Chapter 2


Let's simulate a fair six-sided die roll. Recall that we need to use the `set.seed` function so that we all generate the same random data. For consistency, I always set the seed to 2020, i.e., `set.seed(2020)`. The sample function will generate 100 random die rolls with each possible outcome (1,2,3,4,5,6) being equally likely.

```
set.seed(2020)
data = sample(x = 1:6,size = 100,replace = T); data
```

```
## [1] 4 4 6 1 1 4 2 6 1 5 2 2 6 5 2 3 2 5 4 2 6 6 4 6 4 2 4 5 4 4 3 6 2 2 6 3 5 4 5 5 2 5 1
## [52] 3 2 6 2 1 3 2 1 6 5 5 2 5 6 4 3 6 4 4 2 2 1 6 1 4 6 2 2 2 1 4 1 1 4 2 2 5 5 4 4 4 6 5
```

Let's make a histogram of the simulated data. Recall that each outcome (1,2,3,4,5,6) is equally likely in theory but in our simulation, some numbers may come up more than others. To make the histogram look better, I define the boundary breaks to be 0:6 (i.e., 0,1,2,3,4,5,6).

```
hist(data,breaks=0:6)
```




With only 100 random die rolls, the distribution isn't as flat, uniform, and symmetric as we would expect. The number 3 came up less than we would anticipate. This is OK.

We can also generate a frequency table and use this to create a barplot. This looks a little nicer than the standard histogram and doesn't require use to use the **breaks** argument.

```
table(data)
```

```
## data  
##  1  2  3  4  5  6  
## 15 21  8 21 18 17
```

```
barplot(table(data))
```



If we want a better representation of the true distribution of the fair six-sided die rolls, we could simulate 10000 die rolls instead of 100.


```
set.seed(2020)
```

```
data = sample(x = 1:6, size = 10000, replace = T)
```

```
table(data)
```

```
## data  
##   1    2    3    4    5    6  
## 1630 1657 1733 1655 1629 1696
```

```
barplot(table(data))
```



Simulation can be helpful for us to verify hand calculations and to estimate quantities that are hard to calculate by hand. Our simulation results won't exactly match our hand calculations, but if the sample size is large enough, they should be close.\

For instance, we know that $P(\text{Roll}=1) = 1/6$ for the population of all fair six-sided die rolls. If you refer to the table at the bottom of the previous page, you see that 1 came up 1630 times out of 10000 die rolls. In the simulation, $P(1) = 1630/10000 = 0.1630$ which is pretty close the the theoretical value of $P(1) = 1/6 = 0.1667$.\

Here is an example of how to calculate the probability of a die roll less than 3 in our simulation.

```
# Calculate P(Roll < 3)
# From theory, we know that P(Roll < 3) = P(1) + P(2) = 1/6 + 1/6 = 1/3

# count up the number of rolls less than 3
num.under3 = sum( data < 3 )

# divide the number of rolls less than 3 by the number of rolls
prob.under3 = num.under3 / 10000
prob.under3
```

```
## [1] 0.3287
```

```
# alternatively:
prob.under3 = mean(data < 3)
prob.under3
```

```
## [1] 0.3287
```

From theory, we expect $P(\text{Roll} < 3) = 1/3 = 0.3333$. In our simulation, we found $P(\text{Roll} < 3) = 0.3287$.

Let's simulate rolling two fair six-sided dice 10,000 times and calculate some probabilities. **d1** will represent the first die roll and **d2** will represent the second die roll.

```
set.seed(2020)
d1 = sample(x = 1:6,size = 10000,replace = T)
d2 = sample(x = 1:6,size = 10000,replace = T)
```

From our simulated data, let's calculate $P(D1 = 1 \text{ and } D2 = 2)$, that is, the probability that first die roll is 1 and the second die roll is 2. From theory, this probability should be $1/36 = 0.0278$. (We use the **&** for **AND** in R.)

```
num.events = sum(d1 == 1 & d2 == 2)
prob = num.events / 10000; prob
```

```
## [1] 0.0256
```

Next, let's calculate $P(D1 < 2 \text{ or } D2 > 2)$, that is, the probability that the first die roll is less than 2 or the second die roll is greater than 2. From theory, this probability should be 0.7222. We use the **|** for the **OR** in R.)

```
num.events = sum(d1 < 2 | d2 > 2)
prob = num.events / 10000; prob
```


```
## [1] 0.7183
```

Finally, let's look at the frequency table and distribution of the sum of the two die rolls.

```
table(d1+d2)
```

```
##
##      2      3      4      5      6      7      8      9     10     11     12
##  271   540   824  1154  1373  1703  1360  1149   814   525   287
```

```
barplot(table(d1+d2),xlab="Sum of Dice",ylab="Counts")
```



Discrete Random Variables

In this chapter, we'll explore random variables and their associated probability distributions, with a focus on discrete random variables. These concepts are the backbone of probability theory, enabling us to model and analyze uncertainty. We'll also briefly cover the differences between discrete and continuous random variables—discrete variables take on specific, countable values, while continuous variables can assume any value within a range.

Random Variables

For a given sample space Ω of an experiment, a **random variable (RV)** is a rule that assigns a number to each outcome in Ω . Mathematically, a random variable is a function with the sample space as its domain and the set of real numbers as its range.

A **discrete random variable** is a random variable whose possible values either constitute a finite set or a countably infinite set.

A **continuous** random variable is a random variable whose possible values are uncountably infinite or defined over an interval. Note: $P(X = c) = 0$ for a continuous random variable X and a constant c .

For the following cases, determine whether the random variable is discrete or continuous and state the sample space, Ω .

- (a) Roll a fair six-sided die. Let X be the number on the top of the die.

- (b) Pick a random person and let Y be the number of keys in their pocket.

- (c) Suppose a student is given 50 minutes to complete an exam. Let Z be the amount of time it takes the student to complete the exam.

Probability Distributions for Discrete Random Variables

The **probability mass function (pmf)** or **probability distribution** of a discrete random variable X is defined for every number x by $p(x) = P(X = x)$.

Note: $\sum_{x \in X} p(x) = 1$

The **cumulative distribution function (CDF)** $F(x)$ of a discrete random variable X with pmf $p(x)$ is defined for every number x by: $F(x) = P(X \leq x) = \sum_{y: y \leq x} p(y)$


- $F(x)$ gives you the probability that a random variable will be at most equal to x .
- $F(x)$ must be right-continuous and non-decreasing.
- $F(x)$ must satisfy two limit laws:
 - (i) $\lim_{x \rightarrow -\infty} F(x) = 0$
 - (ii) $\lim_{x \rightarrow \infty} F(x) = 1$

For any two numbers a and b with $a \leq b$, $P(a \leq X \leq b) = F(b) - F(a^-)$.

Roll a fair six-sided die. Let X be the number of the top of the die. Find the pmf and cdf of X and sketch them.



Flip two fair coins. Let Y be the number of coins that land on heads. Find the pmf and cdf of Y and sketch both.



Expected Values (Means)

Let X be a discrete random variable with set of possible outcomes D and pmf $p(x)$. The **expected value** or **mean value** of X , denoted $E[X]$ or μ , is:

$$E[X] = \sum_{x \in D} x \cdot p(x)$$

$$E[h(X)] = \sum p(x) \cdot h(x)$$

$$E[aX + b] = a \cdot E[X] + b$$

Let X have a pmf $p(x)$ and expected value μ . Then the **variance** of X , denoted by $Var(X)$ or σ^2 , is:

$$Var(X) = \sigma^2 = \sum_{x \in D} (x - \mu)^2 \cdot p(x) = E[(X - \mu)^2]$$

The **standard deviation**, denoted SD or σ is:

$$SD(X) = \sigma = \sqrt{Var(X)}$$

$$Var(X) = E[X^2] - (E[X])^2$$

$$Var(aX + b) = a^2 \cdot \sigma^2$$

Roll a fair six-sided die. Let X be the number of the top of the die. Find $E[X]$ and $Var(X)$.

x	1	2	3	4	5	6	Σ
$p(x)$							
$x * p(x)$							
$x^2 * p(x)$							

Flip two fair coins. Let Y be the number of coins that land on heads. Find $E[Y]$ and $Var(Y)$.

x	0	1	2	Σ
$p(x)$				
$x * p(x)$				
$x^2 * p(x)$				

Suppose $Z = 3 \cdot Y + 2$, where Y is defined above. Find $E[Z]$ and $Var(Z)$.

Bernoulli and Binomial Random Variables

Suppose $p(x)$ depends on a quantity that can be assigned any one of a number of possible values, with each different value determining a different probability distribution. Such a quantity is called a **parameter** of the distribution.

The collection of all probability distributions for different values of the parameter is called a **family** of probability distributions.

A **Bernoulli**(p) random variable X is a discrete random variable with two possible outcomes (typically, these outcomes are 0 and 1).

The PMF of a Bernoulli(p) is:

$$p(x) = \begin{cases} 1 - p & x = 0 \\ p & x = 1 \end{cases}$$

If $X \sim \text{Bernoulli}(p)$ random variable, then $E[X] = p$ and $\text{Var}(X) = p(1 - p)$

Show that if X is a Bernoulli(p) random variable, then $E[X] = p$ and $\text{Var}(X) = p(1 - p)$.

X is a **Binomial**(n, p) random variable if X is a discrete random variable that satisfies the following conditions:

- The experiment consists of n Bernoulli trials, where n is fixed.
- The trials are independent.
- The probability of success, p , is constant from trial to trial.
- X is the number of successes

The PMF of a Binomial(n, p) is:


$$p(x) = \binom{n}{x} p^x (1 - p)^{n-x}, x = 0, 1, 2, \dots, n$$

If $X \sim \text{Binomial}(n, p)$ random variable, then $E[X] = np$ and $\text{Var}(X) = np(1 - p)$.

Basketball player LeBron James has a career free throw percentage of 73.1% (i.e., there's a 73.1% chance he will make a basket from the free throw line). Suppose LeBron has six free throw attempts in a game and assume all free throw shots are independent. Answer the following questions.

- (a) Let X be the number of free throws that LeBron makes. Find and sketch the pmf of X .

x	0	1	2	3	4	5	6
$p(x)$							



- (a) On average, how many free throws is LeBron expected to make? What is the variance?

- (c) What is the probability that LeBron makes at least five shots?

Geometric, Discrete Uniform, and Poisson Random Variables

X is a **Geometric**(p) random variable is X is a discrete random variable with the following properties.

- The experiment consists of independent trials.
- Each trial can result in a success or a failure.
- The probability of success is p and is constant across all trials.
- The experiment continues until a successful trial is observed.
- X is the number of total trials

The PMF of a Geometric(p) is:

$$p(x) = p(1 - p)^{x-1}, x = 1, 2, 3, \dots$$

If $X \sim \text{Geometric}(p)$ random variable, then $E[X] = \frac{1}{p}$ and $Var(X) = \frac{1-p}{p^2}$.

Basketball player LeBron James has a career free throw percentage of 73.1%. Suppose he takes shots until makes one. Calculate the probability that 3 or more shots will be required to make his first shot.

X is a **Discrete Uniform(a, b)** random variable is X is a discrete random variable such that the outcomes $a, a + 1, \dots, b$ are equally likely. Let $n = b - a + 1$.

The PMF of a Uniform(a, b) is:

$$p(x) = \frac{1}{n}, x = a, a + 1, \dots, b$$

If $X \sim \text{Uniform}(a, b)$ random variable, then $E[X] = \frac{a+b}{2}$ and $\text{Var}(X) = \frac{(b-a+1)^2-1}{12}$.

Using the formulas above, calculate the expected value and variance of a fair six-sided die roll.

A **Poisson random variable** is a discrete random variable with parameter λ and the following pmf:

$$p(x) = \frac{e^{-\lambda} \lambda^x}{x!}, x = 0, 1, 2, \dots$$

If $X \sim \text{Poisson}(\lambda)$ random variable, then $E[X] = \lambda$ and $\text{Var}(X) = \lambda$.

Note: Poisson random variables are often used to model the number of events that occur in a finite period of time.

Jihan is going fishing and on average, she catches 2 fish per day. Assume that the time between successive fish caught is independent. What is the probability that Jihan catches less than 3 fish in a given day?

R Companion for Chapter 3

We can calculate probabilities for common discrete random variables easily in R. There are two parts to finding probabilities using the PMF or CDF.

R uses four prefixes to reference different elements of a random variable. These are:

- ***p*** for “probability”, the cumulative distribution function (CDF)
- ***q*** for “quantile”, the inverse CDF
- ***d*** for “density”, the probability mass function (PMF)
- ***r*** for “random”, a random variable having the specified distribution

In addition, we have suffixes for common random variables: ***binom*** (binomial), ***pois*** (Poisson), and ***geom*** (Geometric). Note that ***geom*** is defined in terms of failures instead of total trials.

For instance, suppose that $X \sim \text{Binomial}(n = 10, p = 0.5)$, that is, X is a binomial random variable that counts the number of heads in 10 fair coin flips. We can calculate $P(X=5)$, the probability of exactly 5 heads in 10 fair coin flips, as follows.

```
dbinom(x=5,size=10,prob=0.5)
```

```
## [1] 0.2460938
```

If we want the probability of at most five heads in 10 coin flips, $P(X \leq 5)$, we can use `***pbinom`.

```
pbinom(q=5,size=10,prob=0.5)
```

```
## [1] 0.6230469
```

Now, suppose that $Y \sim \text{Poisson}(\lambda = 2)$ and we want to know $P(Y > 3)$. First, we note that $P(Y > 3) = 1 - P(Y \leq 3)$.

```
1-ppois(q=3,lambda=2)
```

```
## [1] 0.1428765
```

If we use the prefix ***r***, we can generate simulated data according to a particular random variable. Let's generate 10,000 observations from a Poisson distribution with $\lambda = 3$, look at the first few values, and then calculate some summary statistics. This will require us to use the ***rpois*** function to generate random Poisson data.

```
rand.data = rpois(n=10000,lambda=3)
head(rand.data) # look at the first few values
```

```
## [1] 3 6 2 4 1 4
```

```
mean(rand.data) # calculate the mean
```

```
## [1] 3.0088
```

```
var(rand.data) # calculate the variance
```

```
## [1] 3.047827
```


From theory, we know that if $Y \sim \text{Poisson}(\lambda = 3)$, then $E[Y] = \text{Var}(Y) = \lambda = 3$. Our simulated results are pretty close to the theoretical values.

```
table(rand.data)
```

```
## rand.data
```

```
##    0    1    2    3    4    5    6    7    8    9   10   11  
##  520 1458 2252 2213 1676 1031  514  211   79   28   15    3
```

```
barplot(table(rand.data))
```




How would you describe the distribution above? It seems to be asymmetric, skewed right (tail on the right), and unimodal.

Let's simulate some data and calculate the probabilities in the sample. Suppose we flip 10 fair coins and we count the number of heads. Let X be the number of heads. We can model X as a binomial random variable with parameters $n=10$ (number of coin flips) and $p=0.5$ (fair coin). Let's simulate running this experiment 100 times.

```
set.seed(2020)
coins = rbinom(n=100,size=10,prob=0.5)
table(coins)
```

```
## coins
##  1  2  3  4  5  6  7  8  9
##  3  1 12 19 32 19  8  5  1
```

```
barplot(table(coins))
```



In our simulation, what was the probability that 5 heads occurred out of 10 flips? In the above table, we can see we got 5 heads on 32 of the 100 simulations, so $P(5) = 32/100 = 0.32$.

We could also calculate as follows:

```
p5 = sum(coins == 5)/100; p5
```

```
## [1] 0.32
```

What is the probability that we got more than 7 heads?

```
p = sum(coins > 7)/100; p
```

```
## [1] 0.06
```